

LEMoE: Advanced Mixture of Experts Adaptor for Lifelong Model Editing of Large Language Models

Renzhi Wang, Piji Li*

College of Computer Science and Technology,

Nanjing University of Aeronautics and Astronautics, China

MIIT Key Laboratory of Pattern Analysis and Machine Intelligence, Nanjing, China

{rzhwang, pjli}@nuaa.edu.cn

Abstract

Large language models (LLMs) require continual knowledge updates to stay abreast of the ever-changing world facts, prompting the formulation of lifelong model editing task. While recent years have witnessed the development of various techniques for single and batch editing, these methods either fail to apply or perform sub-optimally when faced with lifelong editing. In this paper, we introduce LEMoE, an advanced Mixture of Experts (MoE) adaptor for lifelong model editing. We first analyze the factors influencing the effectiveness of conventional MoE adaptor in lifelong editing, including catastrophic forgetting, inconsistent routing and order sensitivity. Based on these insights, we propose a tailored module insertion method to achieve lifelong editing, incorporating a novel KV anchor routing to enhance routing consistency between training and inference stage, along with a concise yet effective clustering-based editing order planning. Experimental results demonstrate the effectiveness of our method in lifelong editing, surpassing previous model editing techniques while maintaining outstanding performance in batch editing task. Our code can be found at: <https://github.com/rzhwang/LEMoE>.

1 Introduction

Large language models (OpenAI, 2023; Touvron et al., 2023a,b; Jiang et al., 2023; Bai et al., 2023) encode a vast amount of world knowledge during pre-training, which can be accessed and utilized through natural language prompts (Petroni et al., 2019). However, the dynamic nature of the real world necessitates regular and continual updates to these models to correct outdated information or integrate new knowledge (Yao et al., 2024; Wang et al., 2024). Also, retraining or fine-tuning of LLMs is often resource-intensive and




Figure 1: The conceptual framework for LEMoE. We align the expert networks in MoE architecture with data batches and freeze the expert networks corresponding to previous data when conducting current edits. $Data_i$ and FFN_i represent the current data and module, with dashed line parts indicating future edits.

time-consuming (Li et al., 2024a), making it impractical for lifelong growing knowledge. Therefore, lifelong model editing (Hartvigsen et al., 2023) has been proposed to remedy the continual knowledge updates and injections for LLMs in a cheap and timely manner (Wang et al., 2024).

In recent years, there has been a proliferation of effective model editing techniques proposed for single or batch editing, such as MEND (Mitchell et al., 2022a), ROME (Meng et al., 2022), MEMIT (Meng et al., 2023), and MEMoE (Wang and Li, 2024). However, these methods often prove inapplicable or exhibit suboptimal performance when faced with lifelong editing task (Wang et al., 2024). In this paper, we introduce LEMoE, an advanced Mixture of Experts (MoE) adaptor, to address the challenges inherent in lifelong editing.

Initially, we analyze the factors that influence the effectiveness of conventional MoE adaptor in lifelong editing, including catastrophic forgetting, inconsistent routing, and order sensitivity. In the Catastrophic Forgetting Analysis (§3.1), we evaluate the performance of conventional MoE adaptor at different positions within the editing sequence to quantify the impact of subsequent edits on preced-

*Corresponding author

ing ones. We observe that the conventional MoE adaptor exhibits significant catastrophic forgetting, where earlier edits are more prone to errors. In the Routing Consistency Analysis (§3.2), we compare the expert networks selected by routing strategy during the training and inference stages when faced with the same input. This comparison reveals a routing inconsistency in the conventional routing strategy, where identical inputs are routed to different experts at different stages. Finally, in the Order Sensitivity Analysis, we highlight that editing order profoundly impacts model performance (§3.3). Through varying the sequence order of the same dataset during lifelong editing, we observe performance variations of up to 20 points, surpassing the improvement of some optimization algorithms.

Based on these insights, we propose a tailored module insertion method to achieve lifelong editing (§4.1). As illustrated in Figure 1, we align the expert networks in the MoE architecture with the data batches in the sequential editing process. When conducting current editing, we freeze the expert network corresponding to the previous data, thereby mitigating the adverse effects of current data editing on previous edits and alleviating catastrophic forgetting from a model mechanism perspective. Secondly, we introduce a novel Key-Value (KV) anchor routing (§4.2), wherein each expert is assigned a key vector and the input instance-level embedding serves as the corresponding value. Based on these key-value pairs, we align the routing computation processes during both training and inference stage. This ensures that identical inputs undergo same routing computation to reach the same expert across all stages, thereby enhancing routing consistency and further mitigating catastrophic forgetting. Finally, leveraging the consistency between the MoE preferences of editing order and the objectives of clustering algorithm, we employ a concise yet effective clustering-based order planning to enhance the overall performance of LEMoE (§4.3).

We conduct experiments on the LLaMA-7B and Mistral-7B models using the ZsRE (Levy et al., 2017) and SelfCheckGPT (Manakul et al., 2023) datasets to evaluate the performance of LEMoE. Experimental results show that our approach surpasses previous model editing methods, while maintaining excellent performance in batch editing.

The main contributions of our work can be summarized as follows:

- We analyze the influential factors of conven-

tional MoE adaptor in lifelong editing task, including catastrophic forgetting, inconsistent routing, and order sensitivity.

- We introduce LEMoE, an advanced MoE adaptor for lifelong model editing. To address the aforementioned challenges, we propose a module insertion method, KV anchor routing, and clustering-based order planning.
- Experimental results show the efficacy of our proposed method in lifelong editing, while simultaneously preserving outstanding performance in batch editing.

2 Preliminaries of Model Editing

Based on previous research (Yao et al., 2023; Zhang et al., 2024; Li et al., 2024a), model editing involves the process of transforming an initial base model f_θ (where θ denotes the model’s parameters) into an edited model $f_{\theta'}$. The goal is to modify the model’s outputs for a specific set of editing instances, while maintaining consistent behavior for all other instances (Li et al., 2024a). The target editing instance can be described as (x_i^e, y_i^e) , with the condition that $f_\theta(x_i^e) \neq y_i^e$. The set of this instances is termed the editing scope I_{edit} , whereas the out-of-scope set O_{edit} comprises inputs not associated with the editing examples. Formally, the criteria for a successful edit can be described as:

$$f_{\theta'}(x_i) = \begin{cases} y_i^e & \text{if } x_i \in I_{edit} \\ f_\theta(x_i) & \text{if } x_i \in O_{edit} \end{cases} \quad (1)$$

We divide model editing tasks into two categories: Batch Editing and Lifelong Editing:

1) **Batch Editing** refers to the simultaneous modification of model f_θ using multiple input instances in one editing operation:

$$\theta' \leftarrow \underset{\theta}{\operatorname{argmin}} \sum_{i=1}^n (\|f_\theta(x_i^e) - y_i^e\|) \quad (2)$$

where n represents the batch size. Batch editing with batch size of 1 is also known as Single Editing.

2) **Lifelong Editing** refers to the continuous iterative modification of model f_θ , also known as Sequential Batch Editing. Lifelong editing use dataset $\mathcal{D}_{edit} = \{\mathcal{B}_1, \mathcal{B}_2, \dots, \mathcal{B}_s\}$ with s sequential batches and each batch $\mathcal{B}_i$ contains n edits:

$$\theta' \leftarrow \underset{\theta}{\operatorname{argmin}} \sum_{j=1}^s \sum_{i=(j-1) \times n + 1}^{j \times n} (\|f_\theta(x_i^e) - y_i^e\|) \quad (3)$$




Figure 2: *Left*: Reliability of conventional MoE under different stage evaluation. “Immediate evaluation” occurs immediately after each edit, “Final evaluation” occurs after all edits in lifelong editing. *Right*: Visualization of routing consistency. The value C_{ij} in each block denotes the proportion of the input data processed by expert i during the training phase that is routed to expert j during the testing phase. Model: LLaMA2-7B. Dataset: ZsRE.

Note that the model is not rolled back to the initial state after each batch editing. Similarly, lifelong editing with batch size of 1 in the sequence is also referred to as Sequential Editing.

Based on the above settings, an effective model editor must satisfy the criteria of three fundamental properties: Reliability, Generalization, and Locality (Yao et al., 2023). These properties are formally defined as follows (Zhang et al., 2024):

1) **Reliability** denotes the average precision of the post-edit model $f_{\theta'}$ concerning the intended edits:

$$\mathbb{E}_{(x_i^e, y_i^e) \sim I_{edit}} \mathbb{1} \{ \arg\max_y f_{\theta'}(y | x_i^e) = y_i^e \} \quad (4)$$

2) **Generalization** quantifies the average precision of the model $f_{\theta'}$ on instances uniformly sampled from the equivalence neighborhood N_{edit} , encompassing input/output pairs pertinent to I_{edit} :

$$\mathbb{E}_{(x_i, y_i^e) \sim N_{edit}} \mathbb{1} \{ \arg\max_y f_{\theta'}(y | x_i) = y_i^e \} \quad (5)$$

3) **Locality** is measured by the proportion at which predictions of the post-edit model $f_{\theta'}$ remain unaltered compared to the pre-edit model f_{θ} :

$$\mathbb{E}_{(x_i, y_i) \sim O_{edit}} \mathbb{1} \{ f_{\theta'}(y | x_i) = f_{\theta}(y | x_i) \} \quad (6)$$

3 Analysis of Influencing Factors

In this section, we analyze the factors that influence the effectiveness of conventional MoE adaptor in lifelong editing, including Catastrophic Forgetting Analysis (§3.1), Routing Consistency Analysis (§3.2), and Order Sensitivity Analysis (§3.3).

3.1 Catastrophic Forgetting Analysis

In the field of continual learning, the general phenomenon of catastrophic forgetting where training

on new tasks degrade performance on old tasks has been extensively reported and studied (Kotha et al., 2023). We aim to investigate whether the MoE adaptor in lifelong model editing also suffers from catastrophic forgetting: whether editing with new data leads to forgetting previously edited data. To assess this, we employ the classic evaluation method for catastrophic forgetting, which involves measuring the performance decrease on previously edited data during the course of lifelong editing.

Experiments To evaluate the impact of current data editing on previous ones during the lifelong editing process, we employ two different stage evaluation methods: (1) a normal evaluation conducted on all editing data only after all edits are completed, and (2) an evaluation conducted immediately after editing the current data to assess the effectiveness of these edits at the current stage. We do not set up another control group where the base model edits only the current data without considering previous data because the accuracy of MoE adaptor under the second evaluation method is nearly 100%. We utilize the LLaMA2-7B as base model and ZsRE dataset (detailed in §5.1). In lifelong editing setting, we perform 100 sequential editing steps, with each step editing a batch of 10 instances, resulting in a total of 1000 edited instances. The evaluation metric is Reliability (detailed in §2). The implementation of conventional MoE adaptor follows (Wang and Li, 2024), employing 4 experts and $top_k = 1$.

Results In the left plot of Figure 2, the reliability of the immediate evaluation during the entire lifelong editing process shows that only once do the score fall below 100. This indicates that the model

consistently achieves desired editing goals at every individual step. However, the model’s overall performance is only around 60 in the final evaluation, with earlier edits exhibiting a more significant decline in effectiveness and some initial edits scoring close to 0. This suggests a pronounced catastrophic forgetting phenomenon, where the model’s forgetfulness of previous editing data markedly diminishes its overall performance in lifelong editing.

3.2 Routing Consistency Analysis

In MoE structure, the specificity of experts directly impacts model performance (Fedus et al., 2022). The design philosophy of MoE adaptor encourages "professional people do professional things", ensuring that the same inputs are routed to the same expert for processing during both training and testing phases (Wang and Li, 2024). We aim to assess the consistency of the routing within the MoE adaptor under lifelong editing setting and explore the degree of specificity among these experts.

Experiments To assess the routing consistency, we log the processing expert for each input during the training phase and compare it against the expert processing the same input during testing. We train models under batch editing setting and lifelong editing setting on identical dataset to compare routing consistency across different tasks. In the lifelong editing setup, sequential editing steps is set to 100 steps, each step editing a batch of 10 instances. Maintaining the same edited data, batch editing utilize a batch size of 1000. We employ LLaMA2-7B and ZsRE, with all other experimental settings consistent with §3.1.

Results On the right side of Figure 2, each sub-graph depicts the proportion C_{ij} where the input processed by expert i during the training phase is routed to expert j during the testing phase. The diagonal element C_{ii} represents the probability that the same input is routed to the same expert during both training and testing phases. Experimental results comparing two editing setups reveal that routing consistency is notably poorer in lifelong editing task, with minimal specificity observed among different experts. In contrast, batch editing exhibits significant routing consistency. Hence, there is a critical need to devise more accurate and effective routing algorithm to guide expert specialization.




Figure 3: *Left*: Performance variability under different editing order. *Right*: Within-Batch/Between-Batch Semantic Similarity v.s. Reliability.

3.3 Order Sensitivity Analysis

In continual learning, the performance of a model significantly varies based on the order of the task arrival sequence (Bell and Lawrence, 2022; Yoon et al., 2020). Previous researches on lifelong editing ignored the impact of editing order on performance. Therefore, we aim to investigate how different editing order affect the overall performance in lifelong editing. Additionally, model tend to learn similar tasks more effectively in continual learning (Bell and Lawrence, 2022), and sentences with high semantic similarity often contain related knowledge. Therefore, we also aim to explore the relationship between the semantic similarity of editing inputs and the editing results.

Experiments To evaluate the model’s editing order sensitivity, we employ the same set of editing data and randomly shuffle the order before performing lifelong editing. In this lifelong editing setup, the sequential editing steps is set to 10, with each step editing a batch of 10 (or 100) instances, resulting in a total of 100 (or 1000) edited instances. Each of these two data volumes experiment is conducted 100 times. To assess the relationship between the semantic similarity of the editing inputs and the editing results, we calculate both within-batch semantic similarity (WBS) and between-batch semantic similarity (BBS) of the editing data. Specially, given dataset $\mathcal{D}_{edit} = \{\mathcal{B}_1, \mathcal{B}_2, \dots, \mathcal{B}_s\}$ with s sequential batches and each batch $\mathcal{B}_i$ contains n edits $\mathcal{B}_i = \{(x_i^e, y_i^e)\}_{i \in [1, n]}$, the WBS_i of $\mathcal{B}_i$ and BBS can be calculated as:

$$WBS_i = \frac{2}{n(n-1)} \sum_{1 \leq i < j \leq n} sim(e_i, e_j) \quad (7)$$

$$BBS = \frac{2}{s(s-1)} \sum_{1 \leq i < j \leq s} sim(\mathcal{B}_i, \mathcal{B}_j) \quad (8)$$




Figure 4: The overall architecture of LEMoE compared with conventional MoE adaptor. We assume that LEMoE is currently at time i to edit data data_i using module FFN_i . *Left*: When editing data data_i , the prior experts corresponding to previous data are all frozen, leaving only the new model FFN_i and router trainable. *Right*: In the training stage, depicted by the solid lines, the routing weight $g(i | x)$ (abbreviated as g_i) is computed using the instance-level embedding and expert key vectors $\{k_1, k_2, \dots, k_i\}$ for expert selection. During inference, as indicated by the dashed lines, the same routing computation is employed to direct the input to the corresponding expert.

where $\text{sim}(\cdot)$ denotes cosine similarity based on embedding, $e_i = \text{concat}(\text{embed}(x_i^e, y_i^e))$ and embed represents the embedding vector output from the model’s embedding layer. For input consisting of multiple tokens, the embedding is the mean for each token. In Equation 8, $B_i = \mathbb{E}_i(e_i)$ denotes the average semantic vector for the i -th group. We employ LLaMA2-7B and ZsRE, with all other experimental settings consistent with §3.1.

Results On the left of Figure 3, the MoE adaptor demonstrates significant order sensitivity. Varying the order of the same editing data leads to performance fluctuations exceeding 20 points. In 100 data editing, these fluctuations even range from 30 to 90 point. This substantial impact of editing order on model performance suggests that adopting order align with model preferences can greatly enhance editing efficiency. The results on the right reveal that higher within-batch semantic similarity and lower between-batch semantic similarity correlate with better editing results. These provide insights for designing more effective editing order.

4 Methods

Based on the above insights, in this section, we provide a detailed introduction to LEMoE an advanced MoE adaptor with new module inserting, KV anchor routing and clustering-based order planning.

4.1 New Module Inserting

Inspired by (Wang and Li, 2024), LEMoE introduces multiple parallel experts within the transformer feed-forward network (FFN) via a bypass mechanism, while freezing all the model’s original parameters. This module is applied in only one transformer block of the entire model. The choice to use the FFN module is motivated not only by its traditional role in MoE but also by recent experimental findings from knowledge probing technologies that suggest the MLP layers within the FFN store knowledge (Geva et al., 2021; Meng et al., 2022, 2023). The bypass mechanism preserves all the original parameters of the model, enhancing the locality of model editing.

However, in conventional MoE adaptor, all experts are sequentially trained without any mechanism to protect prior editing knowledge, which allows current edits to easily affect previous ones and leads to severe catastrophic forgetting. Meanwhile, experimental results indicate that a single FFN expert is sufficient to learn a batch of editing data (Wang and Li, 2024). Therefore, in LEMoE, we adopt a straightforward method to maintain edits from previous learning phases. As shown on the left of Figure 4, when facing a new batch of editing data in the sequence, we add a new FFN module as an expert to learn this batch of data and freeze the expert network corresponding to the previous data. By aligning the expert networks in MoE architecture with the data batches in lifelong editing, we

mitigate the adverse effects of current data editing on previous edits, thereby alleviating catastrophic forgetting from a model mechanism perspective.

Specially, when the $t + 1^{\text{th}}$ batch in lifelong editing dataset arrives, the LEMoE adaptor integrates previous t experts denoted as $f_1, f_2, \dots, f_t$, a router $g(i | x)$ which outputs the corresponding coefficients for each f_i based on the input x along with a newly added expert f_{t+1} . The output h of this module can be expressed as:

$$\begin{aligned} h(x) &= \mathbf{W}_0 \cdot x + \lambda \sum_{i=1}^{t+1} g(i | x) f_i(x) \\ g(i | x) &= \text{Top}_k \left(\frac{e^{r(x)_i}}{\sum_j e^{r(x)_j}} \right) \end{aligned} \quad (9)$$

where $\mathbf{W}_0$ is the frozen original FFN parameters, $r(x)$ is the routing strategy and is modeled by one MLP in conventional MoE. λ is a non-negative weighting coefficient used to balance the old and new knowledge and usually set to 1.

4.2 KV Anchor Routing

We propose the KV anchor routing to align the training and inference processes for expert selection, thereby enhancing routing consistency and addressing catastrophic forgetting at routing level.

During the training phase, when the t -th batch in the lifelong editing dataset arrives, we freeze the parameters of all previous experts $f_1, f_2, \dots, f_{t-1}$ and introduce a new expert f_t to accomplish the current batch editing. We allocate a key vector $\mathbf{k}_i$ for each expert f_i (at time t , only f_t is allocated a new key, while the keys corresponding to previous experts remain frozen) and compute instance-level embedding features of the input as values.

The KV anchor process begins with the j -th input sentence $X_{jt} = \{x_i^{jt}\}_{i=1}^L$ of the current t -th batch data passing through the embedding layer of the LLM backbone to obtain $\mathbf{E}_j^t$ (we omit the superscripts t for simplicity). Since $\mathbf{E}_j \in \mathbb{R}^{m \times d}$ and each key vector $\mathbf{k}_i \in \mathbb{R}^d$ have different sequence lengths, we apply mean-pool operation on the length dimension of $\mathbf{E}_j$, and obtain $\mathbf{e}_j \in \mathbb{R}^d$. Then $\mathbf{e}_j$ is fed to a sub-network to project it into the spaces of the key vectors for better feature alignment. This consists of down and up projection:

$$\mathbf{v}_j = \mathbf{W}^{\text{up}}(\text{SiLU}(\mathbf{W}^{\text{down}} \cdot \mathbf{e}_j)) \quad (10)$$

where $\mathbf{W}^{\text{down}} \in \mathbb{R}^{d_p \times d}$ and $\mathbf{W}^{\text{up}} \in \mathbb{R}^{d \times d_p}$ are learnable projection parameters. Then, the router

$g(i | \mathbf{e}_j)$ in Equation 9 can be defined as:

$$g(i | \mathbf{e}_j) = \text{Top}_k \left(\frac{e^{\mathbf{k}_i \mathbf{v}_j}}{\sum_{i=1}^t e^{\mathbf{k}_i \mathbf{v}_j}} \right) \quad (11)$$

The output for the input token x_i^{jt} of aggregated experts can be obtained:

$$h(x_i^{jt}) = \mathbf{W}_0 \cdot x + \lambda \sum_{i=1}^t g(i | \mathbf{e}_j) f_i(x_i^{jt}) \quad (12)$$

During the inference phase, when testing data from different batches arrive, they undergo the same routing computation of Equations 10 and 11 to reach the appropriate expert. Although there is a possibility that the earlier testing data may be routed to the experts corresponding to the later data, we mitigate this issue by employing the clustering-based editing order selection described in §4.3 which reduces semantic similarity between batches. In summary, aligning the routing computations during training and testing phases through KV anchors enhances routing consistency and further mitigates catastrophic forgetting.

4.3 Clustering-based Order Planning

In §3.3, we observed a correlation between improved editing performance and editing order characterized by high between-batch semantic similarity and low within-batch semantic similarity. This suggests that editing performance can be improved by selecting editing order that align with model biases. Additionally, this objective aligns with the goals of clustering algorithms which aim for high intra-cluster similarity and low inter-cluster similarity. Therefore, we employed the K-means algorithm to group the editing data based on semantic similarity and preferentially selected data from the same cluster for each batch during editing. Experimental results indicate that this straightforward approach is highly effective.

5 Experiments

5.1 Experimental Setups

Datasets and Metrics We used two lifelong model editing datasets: ZsRE (Levy et al., 2017) and SelfCheckGPT (Manakul et al., 2023). ZsRE is a context-free Question Answering (QA) dataset built upon zero-shot relation extraction, and we adopt the split provided by (Zhang et al., 2024). SelfCheckGPT is a dataset for evaluating the performance of model editing methods on mitigating model hallucination, and we followed the

Table 1: Lifelong editing results. **Bold** is the best result. T : Num Edits.

Method	ZsRE								SelfCheckGPT			
	$T = 100$				$T = 1000$				$T = 100$		$T = 600$	
	Rel.↑	Gen.↑	Loc.↑	Avg.↑	Rel.↑	Gen.↑	Loc.↑	Avg.↑	PPL↓	Loc.↑	PPL↓	Loc.↑
LLaMA2-7B												
FT-L	0.30	0.27	0.23	0.27	0.19	0.16	0.03	0.13	33.06	0.41	69.22	0.26
FT-EWC	0.83	0.74	0.08	0.55	0.76	0.69	0.08	0.51	2.10	0.16	4.56	0.24
MEND	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	10.04	0.88	1847.90	0.00
ROME	0.23	0.22	0.04	0.16	0.01	0.01	0.00	0.01	94.15	0.05	104.93	0.02
MEMIT	0.76	0.68	0.85	0.76	0.69	0.65	0.62	0.65	7.18	0.96	13.47	0.94
DEFER	0.20	0.12	0.27	0.20	0.03	0.03	0.74	0.27	8.91	0.19	19.16	0.12
GRACE	0.96	0.00	1.00	0.65	0.97	0.08	1.00	0.68	9.44	1.00	9.34	1.00
MEMoE	0.72	0.46	1.00	0.73	0.70	0.43	1.00	0.71	3.00	1.00	6.59	1.00
LEMoe	0.83	0.62	1.00	0.82	0.80	0.60	1.00	0.80	2.01	1.00	3.36	1.00
Mistral-7B												
FT-L	0.11	0.10	0.02	0.08	0.16	0.13	0.01	0.10	1594.93	0.00	-	-
FT-EWC	0.82	0.72	0.09	0.54	0.76	0.69	0.09	0.51	4.73	0.17	5.46	0.25
MEND	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	23114.94	0.01	-	-
ROME	0.05	0.05	0.02	0.04	0.04	0.04	0.02	0.03	103.75	0.03	241.17	0.01
MEMIT	0.73	0.71	0.88	0.77	0.73	0.70	0.62	0.68	3.22	0.97	7.28	0.95
DEFER	0.28	0.17	0.26	0.24	0.02	0.02	0.67	0.24	9.54	0.43	24.16	0.13
GRACE	1.00	0.00	1.00	0.67	1.00	0.02	1.00	0.67	9.53	1.00	9.57	1.00
MEMoE	0.70	0.43	1.00	0.71	0.70	0.41	1.00	0.70	4.96	1.00	8.91	1.00
LEMoe	0.78	0.52	1.00	0.77	0.75	0.48	1.00	0.74	3.03	1.00	4.39	1.00

GRACE (Hartvigsen et al., 2023) data processing approach. Further details about the datasets are provided in Appendix B.1. In terms of evaluation metrics, we use the three metrics mentioned in §2: Reliability (Rel.), Generalization (Gen.), and Locality (Loc.), along with the average scores (Avg.) over these metrics. Notably, for the SelfCheckGPT dataset, following (Wang et al., 2024), we use the perplexity (PPL) to verify Reliability, and there is no proper metric for generalization.

Baselines We compare LEMoE with the following four types baselines:

- **Fine-tuning based methods:** FT-L (Meng et al., 2022), FT-EWC (Kirkpatrick et al., 2016). FT-L directly fine-tunes a single layer’s FFN and FT-EWC is a continual learning fine-tuning methods based on Elastic Weight Consolidation.
- **Locate and edit methods:** ROME (Meng et al., 2022), MEMIT (Meng et al., 2023). These methods treat FFN of transformer as a linear associative memory apply causal tracing to locate the editing area within model.
- **Meta-learning methods:** MEND (Mitchell et al., 2022a). MEND learns a hyper-network

using additional training data to transform gradient obtained by standard fine-tuning.

- **Memory based methods:** DEFER (Mitchell et al., 2022b), GRACE (Hartvigsen et al., 2023). DEFER is inspired by SERAC (Mitchell et al., 2022b) using an external cache to store explicit editing cases, while GRACE adopts a codebook to store relevant edits.

Implementation Details We selected LLaMA2-7B and Mistral-7B as base models. The modification was applied to layer 18 with $top_k = 1$. Due to limited computational resources, we were able to add a maximum of 5 FFN experts. Consequently, the sequential editing steps were set to 5 and each step contains a batch of 25 (or 200) instances, resulting in a total of 100 (or 1000) editing instances. We use AdamW (Loshchilov and Hutter, 2019) as the optimizer with a learning rate of $2e-4$. Further details are provided in the Appendix B.

5.2 Main Results

Experimental results are presented in Table 1. On ZsRE dataset, LEMoE outperforms all the comparison methods in average scores, achieving up to a 12.68% improvement over the nearest competi-

Table 2: Batch editing results. **Bold** is the best result, and underline is the second-best. ZsRE. LLaMA2-7B.

Method	Rel.↑	Gen.↑	Loc.↑	Avg.↑
FT-L	0.14	0.13	0.70	0.32
MEND	0.01	0.28	0.97	0.34
MEMIT	0.24	0.40	0.17	0.27
SERAC	0.89	0.16	0.81	0.62
GRACE	0.95	0.38	1.00	0.78
MEMoE	1.00	0.90	1.00	0.97
LEMoE	1.00	<u>0.88</u>	1.00	<u>0.96</u>

Table 3: Scaling to 3K edits. ZsRE. LLaMA2-7B.

Method	$T = 2000$				$T = 3000$			
	Rel.	Gen.	Loc.	Avg.	Rel.	Gen.	Loc.	Avg.
GRACE	0.96	0.03	1.00	0.66	0.96	0.03	1.00	0.66
MEMIT	0.64	0.58	0.55	0.59	0.58	0.53	0.47	0.53
LEMoE	<u>0.74</u>	<u>0.50</u>	1.00	0.75	<u>0.70</u>	<u>0.48</u>	1.00	0.73

tor. While MEMIT shows comparable performance at $T=100$, our method demonstrates a substantial performance gap in longer sequence editing task. In Locality, our method consistently scores 1.00, indicating minimal impact on irrelevant inputs. Although GRACE and FT-EWC achieve higher score in Rel. and Gen. respectively, these methods make great sacrifices in Gen. or Loc. Only our method achieves a better balance.

The performance advantage of LEMoE is more pronounced on SelfCheckGPT dataset, maintaining the lowest perplexity score of 3.36 and 4.39 at $T = 600$, with a maximum improvement of 26.31% over the nearest competitor and a constant locality score of 1.00. In summary, across the two datasets and eight baselines, our method shows a clear performance advantage.

6 Detailed Analysis and Discussion

6.1 Batch Editing

Considering the significant performance advantages of conventional MoE adaptor in batch editing (Wang and Li, 2024), we aim to evaluate the changes in batch editing performance of its improved version, LEMoE, after applying the proposed optimizations. The batch size is set to 30 here. As shown in Table 2, LEMoE continues to excel in batch editing, achieving perfect reliability and locality scores of 1.00, with only a slight decline in generalization. Overall, LEMoE’s performance is nearly on par with the original MoE, demonstrating the dual advantages in both batch editing and lifelong editing.

Table 4: Results of ablation study using 1k edits. **Bold** is the best result. ZsRE. LLaMA2-7B.

	Rel.↑	Gen.↑	Loc.↑	Avg.↑
LEMoE	0.82	0.59	1.00	0.80
+ Conventional Routing	0.70	0.43	1.00	0.71
+ Knowledge Routing	0.72	0.48	1.00	0.73
+ Token-level Embed.	0.75	0.46	1.00	0.74
+ Entity-level Embed.	0.80	0.57	1.00	0.79
- Order Planning	0.78	0.55	1.00	0.78
+ Hierarchical Cluster	0.82	0.58	1.00	0.80

6.2 Longer Sequence Editing

We scale the number of lifelong editing to 3K in Table 3. We observe that LEMoE outperforms the strongest baselines MEMIT and GRACE. GRACE excels in reliability but almost entirely loses generalization. While MEMIT demonstrates better generalization, its lower locality scores indicate a significant impact on unrelated data inputs, potentially affecting the model’s general ability (Gu et al., 2024). Only our method achieves a balanced editing performance. Moreover, the performance advantage of our approach increases with the number of edits, highlighting the potential of LEMoE to handle extremely long sequential editing.

6.3 Ablation Study

We present a series of ablation studies to evaluate the influence of various model components, including routing strategies, embedding levels and order planning. The experimental results are shown in Table 4. Conventional routing means the router is modeled by an MLP, knowledge (anchor) routing is the routing strategy in MEMoE and entity-level embedding means substitute the embeddings of named entities from the input for e_j in Equation 10. More details in Appendix B.4.

We observe that: (1) Different model settings exhibit minimal impact on locality but significantly affect generalization. (2) Alteration in routing strategy notably affect reliability and generalization, and conventional routing yields the lowest scores across all metrics. Meanwhile, employing knowledge routing marginally enhances performance yet still lags behind LEMoE, highlighting the pronounced efficacy of KV-anchor routing. (3) Using token-level embeddings for routing inputs notably diminishes model generalization. A possible reason is that token representation may not be suitable for measuring semantic similarity in autoregressive LLMs (Wang et al., 2024), thereby hindering router’s ability to router the same input to

the same expert. (4) Substituting hierarchical clustering for K-means in editing order planning minimally impacts model performance, yet K-means demonstrates higher computational efficiency. This may stem from our utilization of a small number of clusters and a large batch size during dataset construction, which provides clustering algorithms with greater fault tolerance, thereby partially masking the performance differences between the two clustering algorithms.

7 Conclusion

In this paper, We propose LEMoE, an advanced MoE adaptor for lifelong model editing. We analyze three factors influencing the effectiveness of MoE adaptor in lifelong model editing. Then, we propose three optimization modules. These modules align the routing computation processes between training and testing phases, ensuring the same inputs are routed to the same experts. Experimental results validate the effectiveness of LEMoE across multiple models and datasets.

Ethics Statement

Our research on model editing and the proposed LEMoE module adheres to the ethical guidelines outlined by the ACL Ethics Policy. The primary objective of our work is to enhance lifelong editing performance in LLMs. We recognize the critical importance of addressing privacy concerns when model editing publicly accessible, centralized LLMs with private data. And, we acknowledge the potential risks associated with direct parameter edits within models, especially when using harmful data, which require careful mitigation. It's essential to bear in mind that ill-intentioned model editing could lead the model to generate harmful or inappropriate outputs. Therefore, ensuring safe and responsible practices in model editing is of paramount importance. The application of these techniques should be guided by ethical considerations, with safeguards in place to prevent misuse and the production of harmful results. Our commitment to accountability, responsible governance, and continuous ethical assessment underscores our dedication to upholding the highest standards of integrity in the development and deployment of model editing methods.

Limitations

There are several limitations to consider for future directions of model editing of large language models. Firstly, when the learning sequence scales to more data, such as hundreds of batches or tens of thousands of editing instances, continually allocating an expert block for each batch would lead to significant computational and storage costs. Therefore, exploring methods to prune and merge similar experts in the continual learning process presents an interesting research direction. Secondly, our work primarily focuses on the acquisition of factual knowledge, neglecting other types of knowledge. We prioritize the accuracy of knowledge learning while paying less attention to other aspects, such as knowledge reasoning abilities. Thirdly, due to hardware constraints, our investigation was limited to models with up to 7 billion parameters with 5 experts. Additionally, we concentrated on decoder-only autoregressive models, excluding encoder-decoder architectures. Further research that replicates our study using larger-scale models with much more experts and different architecture would be beneficial in confirming our findings.

Acknowledgments

This research is supported by the National Natural Science Foundation of China (No.62476127, No.62106105), the Natural Science Foundation of Jiangsu Province (No.BK20242039), the CCF-Baidu Open Fund (No.CCF-Baidu202307), the CCF-Zhipu AI Large Model Fund (No.CCF-Zhipu202315), the Fundamental Research Funds for the Central Universities (No.NJ2023032), the Scientific Research Starting Foundation of Nanjing University of Aeronautics and Astronautics (No.YQR21022), and the High Performance Computing Platform of Nanjing University of Aeronautics and Astronautics.

References

- Rahaf Aljundi, Eugene Belilovsky, Tinne Tuytelaars, Laurent Charlin, Massimo Caccia, Min Lin, and Lucas Page-Caccia. 2019. [Online continual learning with maximal interfered retrieval](#). In *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pages 11849–11860.
- Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei

- Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. 2023. [Qwen technical report](#). *CoRR*, abs/2309.16609.
- Samuel J. Bell and Neil D. Lawrence. 2022. [The effect of task ordering in continual learning](#). *CoRR*, abs/2205.13323.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. [Language models are few-shot learners](#). In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*.
- Alexandra Chronopoulou, Matthew E. Peters, Alexander Fraser, and Jesse Dodge. 2023. [Adaptersoup: Weight averaging to improve generalization of pre-trained language models](#). In *Findings of the Association for Computational Linguistics: EACL 2023, Dubrovnik, Croatia, May 2-6, 2023*, pages 2009–2018. Association for Computational Linguistics.
- Together Computer. 2023. [Redpajama: an open dataset for training large language models](#).
- Damai Dai, Li Dong, Yaru Hao, Zhifang Sui, Baobao Chang, and Furu Wei. 2022. [Knowledge neurons in pretrained transformers](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pages 8493–8502. Association for Computational Linguistics.
- Qingxiu Dong, Damai Dai, Yifan Song, Jingjing Xu, Zhifang Sui, and Lei Li. 2022. [Calibrating factual knowledge in pretrained language models](#). In *Findings of the Association for Computational Linguistics: EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 5937–5947. Association for Computational Linguistics.
- Zhibin Duan, Hao Zhang, Chaojie Wang, Zhengjue Wang, Bo Chen, and Mingyuan Zhou. 2021. [Enslm: Ensemble language model for data diversity by semantic clustering](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing, ACL/IJCNLP 2021, (Volume 1: Long Papers), Virtual Event, August 1-6, 2021*, pages 2954–2967. Association for Computational Linguistics.
- William Fedus, Barret Zoph, and Noam Shazeer. 2022. [Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity](#). *J. Mach. Learn. Res.*, 23:120:1–120:39.
- Chris Fifty, Ehsan Amid, Zhe Zhao, Tianhe Yu, Rohan Anil, and Chelsea Finn. 2021. [Efficiently identifying task groupings for multi-task learning](#). In *Advances in Neural Information Processing Systems 34: Annual Conference on Neural Information Processing Systems 2021, NeurIPS 2021, December 6-14, 2021, virtual*, pages 27503–27516.
- Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. 2021. [Transformer feed-forward layers are key-value memories](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, EMNLP 2021, Virtual Event / Punta Cana, Dominican Republic, 7-11 November, 2021*, pages 5484–5495. Association for Computational Linguistics.
- Yunhao Gou, Zhili Liu, Kai Chen, Lanqing Hong, Hang Xu, Aoxue Li, Dit-Yan Yeung, James T. Kwok, and Yu Zhang. 2023. [Mixture of cluster-conditional lora experts for vision-language instruction tuning](#). *CoRR*, abs/2312.12379.
- Jia-Chen Gu, Hao-Xiang Xu, Jun-Yu Ma, Pan Lu, Zhen-Hua Ling, Kai-Wei Chang, and Nanyun Peng. 2024. [Model editing can hurt general abilities of large language models](#). *CoRR*, abs/2401.04700.
- Shashank Gupta, Subhabrata Mukherjee, Krishan Subudhi, Eduardo Gonzalez, Damien Jose, Ahmed Hassan Awadallah, and Jianfeng Gao. 2022. [Sparsely activated mixture-of-experts are robust multi-task learners](#). *CoRR*, abs/2204.07689.
- Suchin Gururangan, Margaret Li, Mike Lewis, Weijia Shi, Tim Althoff, Noah A. Smith, and Luke Zettlemoyer. 2023. [Scaling expert language models with unsupervised domain discovery](#). *CoRR*, abs/2303.14177.
- Tom Hartvigsen, Swami Sankaranarayanan, Hamid Palangi, Yoon Kim, and Marzyeh Ghassemi. 2023. [Aging with GRACE: lifelong model editing with discrete key-value adapters](#). In *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023, New Orleans, LA, USA, December 10 - 16, 2023*.
- Thomas Henn, Yasukazu Sakamoto, Clément Jacquet, Shunsuke Yoshizawa, Masamichi Andou, Stephen Tchen, Ryosuke Saga, Hiroyuki Ishihara, Katsuhiko Shimizu, Yingzhen Li, and Ryutaro Tanno. 2021. [A principled approach to failure analysis and model repairment: Demonstration in medical imaging](#). In *Medical Image Computing and Computer Assisted*

- Intervention - MICCAI 2021 - 24th International Conference, Strasbourg, France, September 27 - October 1, 2021, Proceedings, Part III*, volume 12903 of *Lecture Notes in Computer Science*, pages 509–518. Springer.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. [Lora: Low-rank adaptation of large language models](#). In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net.
- Zeyu Huang, Yikang Shen, Xiaofeng Zhang, Jie Zhou, Wenge Rong, and Zhang Xiong. 2023. [Transformer-patcher: One mistake worth one neuron](#). In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net.
- Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de Las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, L  lio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timoth  e Lacroix, and William El Sayed. 2023. [Mistral 7b](#). *CoRR*, abs/2310.06825.
- Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de Las Casas, Emma Bou Hanna, Florian Bressand, Gianna Lengyel, Guillaume Bour, Guillaume Lample, L  lio Renard Lavaud, Lucile Saulnier, Marie-Anne Lachaux, Pierre Stock, Sandeep Subramanian, Sophia Yang, Szymon Antoniak, Teven Le Scao, Th  ophile Gerv  t, Thibaut Lavril, Thomas Wang, Timoth  e Lacroix, and William El Sayed. 2024. [Mistral of experts](#). *CoRR*, abs/2401.04088.
- James Kirkpatrick, Razvan Pascanu, Neil C. Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. 2016. [Overcoming catastrophic forgetting in neural networks](#). *CoRR*, abs/1612.00796.
- Suhas Kotha, Jacob Mitchell Springer, and Aditi Raghunathan. 2023. [Understanding catastrophic forgetting in language models via implicit inference](#). *CoRR*, abs/2309.10105.
- Matthias De Lange, Rahaf Aljundi, Marc Masana, Sarah Parisot, Xu Jia, Ales Leonardis, Gregory G. Slabaugh, and Tinne Tuytelaars. 2022. [A continual learning survey: Defying forgetting in classification tasks](#). *IEEE Trans. Pattern Anal. Mach. Intell.*, 44(7):3366–3385.
- Omer Levy, Minjoon Seo, Eunsol Choi, and Luke Zettlemoyer. 2017. [Zero-shot relation extraction via reading comprehension](#). In *Proceedings of the 21st Conference on Computational Natural Language Learning (CoNLL 2017), Vancouver, Canada, August 3-4, 2017*, pages 333–342. Association for Computational Linguistics.
- Daliang Li, Ankit Singh Rawat, Manzil Zaheer, Xin Wang, Michal Lukasik, Andreas Veit, Felix X. Yu, and Sanjiv Kumar. 2023a. [Large language models with controllable working memory](#). In *Findings of the Association for Computational Linguistics: ACL 2023, Toronto, Canada, July 9-14, 2023*, pages 1774–1793. Association for Computational Linguistics.
- Shuaiyi Li, Yang Deng, Deng Cai, Hongyuan Lu, Liang Chen, and Wai Lam. 2024a. [Consecutive model editing with batch alongside hook layers](#). *CoRR*, abs/2403.05330.
- Shuaiyi Li, Yang Deng, Deng Cai, Hongyuan Lu, Liang Chen, and Wai Lam. 2024b. [Consecutive model editing with batch alongside hook layers](#). *CoRR*, abs/2403.05330.
- Xiaopeng Li, Shasha Li, Shezheng Song, Jing Yang, Jun Ma, and Jie Yu. 2023b. [PMET: precise model editing in a transformer](#). *CoRR*, abs/2308.08742.
- Bill Yuchen Lin, Sida Wang, Xi Victoria Lin, Robin Jia, Lin Xiao, Xiang Ren, and Scott Yih. 2022. [On continual model refinement in out-of-distribution data streams](#). In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pages 3128–3139. Association for Computational Linguistics.
- Zhenhua Liu, Yunhe Wang, Kai Han, Wei Zhang, Siwei Ma, and Wen Gao. 2021. [Post-training quantization for vision transformer](#). In *Advances in Neural Information Processing Systems 34: Annual Conference on Neural Information Processing Systems 2021, NeurIPS 2021, December 6-14, 2021, virtual*, pages 28092–28103.
- Ilya Loshchilov and Frank Hutter. 2019. [Decoupled weight decay regularization](#). In *7th International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*. OpenReview.net.
- Keming Lu, Hongyi Yuan, Runji Lin, Junyang Lin, Zheng Yuan, Chang Zhou, and Jingren Zhou. 2023. [Routing to the expert: Efficient reward-guided ensemble of large language models](#). *CoRR*, abs/2311.08692.
- Aman Madaan, Niket Tandon, Peter Clark, and Yiming Yang. 2022. [Memory-assisted prompt editing to improve GPT-3 after deployment](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 2833–2861. Association for Computational Linguistics.
- Potsawee Manakul, Adian Liusie, and Mark J. F. Gales. 2023. [Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language](#)

- models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6-10, 2023*, pages 9004–9017. Association for Computational Linguistics.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. 2022. [Locating and editing factual associations in GPT](#). In *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*.
- Kevin Meng, Arnab Sen Sharma, Alex J. Andonian, Yonatan Belinkov, and David Bau. 2023. [Mass-editing memory in a transformer](#). In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net.
- Eric Mitchell, Charles Lin, Antoine Bosselut, Chelsea Finn, and Christopher D. Manning. 2022a. [Fast model editing at scale](#). In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net.
- Eric Mitchell, Charles Lin, Antoine Bosselut, Christopher D. Manning, and Chelsea Finn. 2022b. [Memory-based model editing at scale](#). In *International Conference on Machine Learning, ICML 2022, 17-23 July 2022, Baltimore, Maryland, USA*, volume 162 of *Proceedings of Machine Learning Research*, pages 15817–15831. PMLR.
- Shikhar Murty, Christopher D. Manning, Scott M. Lundberg, and Marco Túlio Ribeiro. 2022. [Fixing model bugs with natural language patches](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 11600–11613. Association for Computational Linguistics.
- OpenAI. 2023. [GPT-4 technical report](#). *CoRR*, abs/2303.08774.
- Fabio Petroni, Tim Rocktäschel, Sebastian Riedel, Patrick S. H. Lewis, Anton Bakhtin, Yuxiang Wu, and Alexander H. Miller. 2019. [Language models as knowledge bases?](#) In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, EMNLP-IJCNLP 2019, Hong Kong, China, November 3-7, 2019*, pages 2463–2473. Association for Computational Linguistics.
- David Rolnick, Arun Ahuja, Jonathan Schwarz, Timothy P. Lillicrap, and Gregory Wayne. 2019. [Experience replay for continual learning](#). In *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pages 348–358.
- Haizhou Shi, Zihao Xu, Hengyi Wang, Weiyi Qin, Wenyuan Wang, Yibin Wang, and Hao Wang. 2024. [Continual learning of large language models: A comprehensive survey](#). *CoRR*, abs/2404.16789.
- Tal Shnitzer, Anthony Ou, Mírian Silva, Kate Soule, Yuekai Sun, Justin Solomon, Neil Thompson, and Mikhail Yurochkin. 2023. [Large language model routing with benchmark datasets](#). *CoRR*, abs/2309.15789.
- Chenmian Tan, Ge Zhang, and Jie Fu. 2023. [Massive editing for large language models via meta learning](#). *CoRR*, abs/2311.04661.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurélien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. 2023a. [Llama: Open and efficient foundation language models](#). *CoRR*, abs/2302.13971.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shrutu Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton-Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian Khabsa, Isabel Kloumann, Artem Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rashi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, Ranjan Subramanian, Xiaoqing Ellen Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zheng Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurélien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. 2023b. [Llama 2: Open foundation and fine-tuned chat models](#). *CoRR*, abs/2307.09288.
- Frederik Träuble, Anirudh Goyal, Nasim Rahaman, Michael Curtis Mozer, Kenji Kawaguchi, Yoshua Bengio, and Bernhard Schölkopf. 2023. [Discrete key-value bottleneck](#). In *International Conference on Machine Learning, ICML 2023, 23-29 July 2023, Honolulu, Hawaii, USA*, volume 202 of *Proceedings of Machine Learning Research*, pages 34431–34455. PMLR.
- Aäron van den Oord, Oriol Vinyals, and Koray Kavukcuoglu. 2017. [Neural discrete representation learning](#). In *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, pages 6306–6315.
- Tu Vu, Tong Wang, Tsendsuren Munkhdalai, Alessandro Sordani, Adam Trischler, Andrew Mattarella-

- Micke, Subhansu Maji, and Mohit Iyyer. 2020. [Exploring and predicting transferability across NLP tasks](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing, EMNLP 2020, Online, November 16-20, 2020*, pages 7882–7926. Association for Computational Linguistics.
- Peng Wang, Zexi Li, Ningyu Zhang, Ziwen Xu, Yunzhi Yao, Yong Jiang, Pengjun Xie, Fei Huang, and Huajun Chen. 2024. [Wise: Rethinking the knowledge memory for lifelong model editing of large language models](#). *arXiv preprint arXiv:2405.14768*.
- Renzhi Wang and Piji Li. 2024. [Memoe: Enhancing model editing with mixture of experts adapters](#). *arXiv preprint arXiv:2405.19086*.
- Yaqing Wang, Sahaj Agarwal, Subhabrata Mukherjee, Xiaodong Liu, Jing Gao, Ahmed Hassan Awadallah, and Jianfeng Gao. 2022. [Adamix: Mixture-of-adaptations for parameter-efficient model tuning](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 5744–5760. Association for Computational Linguistics.
- Tongtong Wu, Linhao Luo, Yuan-Fang Li, Shirui Pan, Thuy-Trang Vu, and Gholamreza Haffari. 2024a. [Continual learning for large language models: A survey](#). *CoRR*, abs/2402.01364.
- Xun Wu, Shaohan Huang, and Furu Wei. 2024b. [Mixture of lora experts](#). *CoRR*, abs/2404.13628.
- Yunzhi Yao, Peng Wang, Bozhong Tian, Siyuan Cheng, Zhoubo Li, Shumin Deng, Huajun Chen, and Ningyu Zhang. 2023. [Editing large language models: Problems, methods, and opportunities](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6-10, 2023*, pages 10222–10240. Association for Computational Linguistics.
- Zihan Yao, Yu He, Tianyu Qi, and Ming Li. 2024. [Scalable model editing via customized expert networks](#). *CoRR*, abs/2404.02699.
- Qinyuan Ye, Juan Zha, and Xiang Ren. 2022. [Eliciting and understanding cross-task skills with task-level mixture-of-experts](#). In *Findings of the Association for Computational Linguistics: EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 2567–2592. Association for Computational Linguistics.
- Jaehong Yoon, Saehoon Kim, Eunho Yang, and Sung Ju Hwang. 2020. [Scalable and order-robust continual learning with additive parameter decomposition](#). In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net.
- Ted Zadouri, Ahmet Üstün, Arash Ahmadian, Beyza Ermis, Acyr Locatelli, and Sara Hooker. 2023. [Pushing mixture of experts to the limit: Extremely parameter efficient moe for instruction tuning](#). *CoRR*, abs/2309.05444.
- Ningyu Zhang, Yunzhi Yao, Bozhong Tian, Peng Wang, Shumin Deng, Mengru Wang, Zekun Xi, Shengyu Mao, Jintian Zhang, Yuansheng Ni, Siyuan Cheng, Ziwen Xu, Xin Xu, Jia-Chen Gu, Yong Jiang, Pengjun Xie, Fei Huang, Lei Liang, Zhiqiang Zhang, Xiaowei Zhu, Jun Zhou, and Huajun Chen. 2024. [A comprehensive study of knowledge editing for large language models](#). *CoRR*, abs/2401.01286.
- Ce Zheng, Lei Li, Qingxiu Dong, Yuxuan Fan, Zhiyong Wu, Jingjing Xu, and Baobao Chang. 2023. [Can we edit factual knowledge by in-context learning?](#) In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6-10, 2023*, pages 4862–4876. Association for Computational Linguistics.
- Wangchunshu Zhou, Canwen Xu, and Julian J. McAuley. 2022. [Efficiently tuned parameters are task embeddings](#). In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pages 5007–5014. Association for Computational Linguistics.

A Related Work

A.1 Model Editing

Model editing is a new and active research area where the goal is to make targeted changes to a pre-trained model’s behavior (Zhang et al., 2024). Given the fast-growing parameter sizes of LLMs, frequently updating LLMs with new knowledge through retraining is more and more expensive. Hence, it is vital to effectively edit the LLMs’ knowledge without retraining. Previous studies have explored multiple methods for editing the knowledge of LLMs, which can be broadly categorized into two streams based on whether it alters the parameters of the original model (Yao et al., 2023; Zhang et al., 2024):

Preserve model parameters: (1) **Retrieve augmentation.** These techniques leverage an external knowledge base to enrich or correct information accessible to language models. These augmented knowledge bases seamlessly integrate with the base model, enabling effective retrieval of relevant information when prompted (Murty et al., 2022; Madaan et al., 2022; Li et al., 2023a). For example, IKE (Zheng et al., 2023) employs an in-context learning approach that adjusts language model outputs using corpus-based demonstrations

guided by similarity metrics, thereby obviating the need for gradient-based adjustments. (2) **Adding additional parameters:** This paradigm involves introducing additional trainable parameters to augment a language model's existing knowledge, while preserving its original parameters in a frozen state. T-Patcher (Huang et al., 2023) and CaliNET (Dong et al., 2022) exemplify this paradigm by integrating specific neurons or patches into the final layer of their Feed-Forward Networks. T-Patcher assigns individual neurons to each distinct error, while CaliNET incorporates multiple neurons to handle various knowledge scenarios. In contrast, GRACE (Hartvigsen et al., 2023) employs a discrete codebook mechanism to dynamically add and update elements, enhancing the model's predictive capabilities over time. (3) **Meta learning** Recent meta-learning methods use hypernetworks for aiding editing. MEND (Mitchell et al., 2022a) introduces a hypernetwork designed to decouple fine-tuning gradients into updates that generalize edits without compromising performance on unrelated inputs. To mitigate the cancellation issue inherent in MEND, MALMEN (Tan et al., 2023) employs a hyper-network to generate weight shifts for editing, formulating the aggregation of these shifts as a least squares problem.

Modify model parameters: This methodology begins by identifying parameters associated with specific knowledge and directly adjusting them. The Knowledge Neuron (KN) approach (Dai et al., 2022) introduces a technique to attribute knowledge to individual "knowledge neurons" and subsequently updates these neurons accordingly. ROME (Meng et al., 2022) utilizes causal mediation analysis to pinpoint areas requiring modification. Both KN and ROME operate under the constraint of editing one factual association at a time. To overcome this limitation, MEMIT (Meng et al., 2023) extends ROME's framework, enabling simultaneous editing across multiple instances. Building on MEMIT, PMET (Li et al., 2023b) integrates attention values to achieve superior performance enhancements. COMEBA-HK (Li et al., 2024b) identifies the Local Editing Scope and extends MEMIT for sequential editing.

A.2 Mixture of Experts

The concept of MoE, particularly when combined with sparse routing, is recognized for significantly enhancing model capacity with minimal computa-

tional overhead (Fedus et al., 2022). Key distinctions in this approach include: i) adapter experts are not trained during the pre-training of the base model, ii) they are parameter-efficient, and iii) they are tailored to specific tasks, unlike token-level opaque computation units whose specialization is not easily interpretable (Jiang et al., 2024). Regarding the second point, (Wang et al., 2022; Zadouri et al., 2023) utilize routing each example to a set of experts, demonstrating improved performance on unseen tasks. (Gupta et al., 2022) implements a separate router for each task and selects a router from a similar task based on domain knowledge. (Ye et al., 2022) proposes task-level MoEs, where a collection of transformer layers acts as experts, and a router dynamically selects from these experts. Additionally, several recent studies have proposed methods for routing queries to specialized pretrained open-source LLMs (Lu et al., 2023; Shnitzer et al., 2023).

A.3 Continual Learning

Continual Learning (CL) (Shi et al., 2024; Wu et al., 2024a) is an essential aspect of machine learning as it enables models to adapt to new tasks while retaining performance on previous ones. It mainly focus on the issue of catastrophic forgetting in deep learning models when exposed to new knowledge (Lange et al., 2022). Recent research has explored diverse approaches in this domain. Among these approaches, continual fine-tuning stands out, involving the iterative refinement of LLMs with incoming instances. For instance, (Lin et al., 2022) conducts an extensive investigation into this method. However, it has been noted that integrating regularized fine-tuning techniques such as Elastic Weight Consolidation (Kirkpatrick et al., 2016), Experience Replay (Rolnick et al., 2019), and Maximally Interfered Replay (Aljundi et al., 2019) can lead to a decline in performance on earlier tasks while preserving some memory of past inputs. This observation underscores the challenges unique to editing in contrast to conventional continual fine-tuning (Henn et al., 2021), particularly given the uneven distribution of edits. One promising avenue in continual learning involves the adoption of key-value methodologies, inspired by advancements in computer vision (Liu et al., 2021; van den Oord et al., 2017). Notably, discrete key-value methods have proven effective in managing shifting distributions (Träuble et al., 2023). These methods cache values to ensure in-

puts remain within distribution bounds for downstream encoders, thereby enabling the integration of longer-term memory, contingent on available computational resources.

A.4 Data Clustering for LLMs

Data clustering methods for LLMs have been proposed to enhance performance and reduce task interference (Fifty et al., 2021; Gururangan et al., 2023; Gou et al., 2023). These methods include clustering based on similarities computed using tf-idf and neural embeddings, K-means clustering with balanced linear assignment, and soft clustering with Gaussian Mixture Models (GMMs) (Chronopoulou et al., 2023; Gururangan et al., 2023; Duan et al., 2021). Recent work by (Zhou et al., 2022) highlights the potential of adapter parameters as effective task embeddings for clustering. Additionally, a similar observation regarding task gradients has been made by (Vu et al., 2020).

B Implementation Details

B.1 Datasets Details

ZsRE The ZsRE dataset is a context-free Question Answering (QA) dataset that has been extensively studied in the model editing literature (Meng et al., 2022, 2023; Mitchell et al., 2022b; Hartvigsen et al., 2023; Wang and Li, 2024). Each record in this dataset includes an editing statement $\mathbf{x}_i^e$ with target answer $\mathbf{y}_i^e$, a paraphrase prompt $\mathbf{x}_{\text{gen}_i}$ and a locality prompt $\mathbf{x}_{\text{loc}}$. We adopt the same train/test split as (Mitchell et al., 2022a), consisting of 163,196 training examples and 19,086 test examples. Notably, MEND is the only method that requires fitting a hyper network on the training set; other methods discard the training set and directly perform edits and evaluations on the test set. For our experiments, we randomly sampled 1k and 3k records from the test set to form the edit sets.

SelfCheckGPT We employ SelfCheckGPT (Manakul et al., 2023), the same dataset as GRACE, to evaluate the effectiveness of Model Editors in reducing hallucinations in autoregressive language models. This dataset consists of highly inaccurate sentences generated from GPT-3 (Brown et al., 2020), which are then replaced with corresponding accurate sentences from Wikipedia. This setup mirrors real-world deployment scenarios where models exhibit "unexpected behaviors". The edits in this dataset are significantly longer compared to ZsRE, presenting a more challenging editing

environment. Unlike GRACE, which utilized GPT2-XL (1.5B), our primary experiments use larger LLMs, specifically LLaMA and Mistral, each with 7B parameters. We measure the retention of $\mathbf{x}_{\text{loc}}$ from the base model, RedPajama (Computer, 2023), a publicly available version of LLaMA’s pre-training data.

B.2 Implementation of Baselines

FT-L We followed the procedures outlined in (Wang et al., 2024): all other layers of the LLMs remain frozen, and only a single MLP layer undergoes fine-tuning using an autoregressive loss function. Furthermore, we impose a L_∞ norm constraint to ensure that the parameters do not deviate significantly from the pretrained distribution. Employ the Adam optimizer with consideration of learning rates at 1e-5, 1e-4, and 5e-4, and conduct gradient descents for 50 iterations, ultimately reporting the best results at a learning rate of 5e-4.

FT-EWC Elastic Weight Consolidation (EWC) effectively mitigates catastrophic forgetting by updating model weights using the Fisher information matrix, which is computed based on past parameter updates and scaled by a factor λ (Kirkpatrick et al., 2016). In line with (Hartvigsen et al., 2023), our implementation does not incorporate L_∞ norm constraints, setting the learning rate at 1e-2, the λ_{ewc} penalty factor at 0.1, and the number of replay instances at 10.

MEND MEND (Mitchell et al., 2022a) performs model editing by employing a hyper-network to transform the gradients derived from standard fine-tuning. This process involves decomposing the model gradients into a low-rank format (rank=1) before converting them into new gradients, which are subsequently applied to the target layer for parameter updates. During training, a small auxiliary hyper-network processes editing examples $(\mathbf{x}_i^e, \mathbf{y}_i^e)$ and $(\mathbf{x}_{\text{gen}_i}, \mathbf{y}_i^e)$. The training loss for MEND consists of the standard autoregressive loss combined with the KL divergence loss, measuring the model’s output on $(\mathbf{x}_{\text{gen}_i}, \mathbf{y}_i^e)$ before and after editing. This hyper-network is pivotal in the editing procedure. Due to the substantial computational resources required to train the meta-network for, the results are from (Wang et al., 2024).

ROME ROME (Meng et al., 2022) employs causal analysis to identify knowledge residing in specific MLP layers and refines the entire matrix

Table 5: An editing dataset example from ZsRE and SelfCheckGPT.

Dataset	Type	Text
ZsRE	$\mathbf{x}_i^e, \mathbf{y}_i^e$	Which continent is Berkner Island in? South America
	$\mathbf{x}_{loc_i}, \mathbf{y}_{loc}$	who gets the golden boot if its a tie? shared
	$\mathbf{x}_{gen_i}, \mathbf{y}_i^e$	On which continent is Berkner Island located? South America
SelfCheckGPT	$\mathbf{x}_i^e, \mathbf{y}_i^e$	This is a Wikipedia passage about heinz christian pander. Heinz Christian Pander (1794 - 1865) was a German anatomist and embryologist who was born in Riga, Latvia. He studied medicine at the University of Dorpat and later at the University of Berlin. In 1820, he took part in a scientific expedition to Bokhara as a naturalist.
	$\mathbf{x}_{loc}, \mathbf{y}_{loc}$	Tired and restlessly, drifting in and out of sleep. Hearing crashing and banging, thinking the roof will cave in. Not alert enough to quite know what. it was, I yelled loudly for whoever was making those noises at such an hour to stop. They heard and listened, I'm guessing

Table 6: Dataset statistics for main results. *Locality Data* is the irrelevant data of the editing process. T is the number of samples. *Pre-edit* is the unedited model’s performance on each dataset.

SETTING	EDITING DATA	T	Pre-edit (LLaMA/Mistral)
QA	ZsRE	1,000	0.36/0.39 ACC
Hallucination	SelfCheckGPT	600	27.4/19.4 PPL

via least squares approximation. This approach assumes MLP as the central repository of knowledge (Geva et al., 2021), incrementally injecting individual pieces of information into the MLP through a Lagrangian residual term at each iteration. Following (Wang et al., 2024), in LLaMA and Mistral, ROME edits the fifth layer, while MEMIT edits layers [4,5,6,7,8].

MEMIT The MEMIT utilized in this study, denoted as MEMIT-MASS as introduced in (Wang et al., 2024), differs notably from its original counterpart. In contrast to sequential editing, MEMIT-MASS facilitates batch processing for modifying multiple knowledge fragments concurrently. Suppose we collect streaming errors as $(\mathcal{X}, \mathcal{Y}) = \{(\mathbf{x}_0, \mathbf{y}_0), (\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_T, \mathbf{y}_T)\}$ and inject them collectively into the MLP, it only involves a single editing operation on the original model as $f_{\Theta_T} = \text{MEMIT}(f_{\Theta_0}, \mathcal{X}, \mathcal{Y})$. Despite its drawback of lacking real-time correction capabilities, we include this approach as a baseline in our experimental evaluations, given the extremely bad performance of the original MEMIT framework.

DEFER In GRACE, a reimplementation of SERAC (Mitchell et al., 2022b) is utilized, denoted as DEFER. DEFER integrates a network denoted as g (akin to the scope classifier in SERAC). This net-

work g predicts whether to rely on: 1) predictions from the LLMs, or 2) predictions from a newly introduced model. This new model, configured as a single-layer linear network o with a sigmoid activation function, parallels the counterfactual model in SERAC. Throughout the editing phase, g and o undergo joint fine-tuning processes. The experiment with learning rates of $7e-5$, $7e-4$, and $1e-3$, and ultimately report using $7e-5$ (optimal).

GRACE GRACE (Hartvigsen et al., 2023) utilizes a discrete key-value codebook and maintains the codebook throughout the editing flow by adding, expanding, and splitting KEYS. During the inference phase, it retrieves the nearest KEY and determines whether to replace the activation of the hidden layer output. We adhere to the meticulously crafted parameters outlined in the original study, configuring the optimization of the learning rate to a value of 1 and using “replace last” to only replace the activation of the last token in autoregressive scenarios.. The iterative process for optimizing these values spans 100 cycles, with an initial $\epsilon = 1$.

MEMoE MEMoE (Wang and Li, 2024) updates knowledge using a bypass MoE structure, keeping the original parameters unchanged to preserve the general ability of LLMs. And, the knowledge anchor routing ensures that inputs requiring similar knowledge are routed to the same expert, thereby enhancing the generalization of the updated knowledge. Following the parameters identified in the original paper, we consulted 4 experts, setting the *top-k* value to 1 and a learning rate of $2e-4$. The modification is applied to `model.layers[16].mlp.up_proj.weight` and `model.layers[16].mlp.down_proj.weight`.

We also adopt auxiliary loss for balancing the top-k selection of routing following (Fedus et al., 2022).

B.3 Training Details of LEMoE

The training loss for the attentive learning of the t -th batch data $\mathcal{B}_t$ is:

$$L_{\text{task}} = - \sum_{(x_t, y_t) \in \mathcal{T}_t} \log P(y_t | x_t; \theta_m, \theta_f, \theta_{\text{proj}}, \theta_k) \quad (13)$$

where $\theta_m, \theta_f, \theta_{\text{proj}}$ and θ_k are parameters of the LLM backbone, the experts, the query projection layer and the set of all key vectors, respectively. And only those parameters belongs to the current t -th task are trainable, including $\theta_{f_t}, \theta_{\text{proj}}$ and θ_{k_t} .

The hyperparameters for the ZsRE and Self-CheckGPT are identical. Specially, We use the AdamW (Loshchilov and Hutter, 2019) as the optimizer with a learning rate of $2e-4$. The modification of the model is applied to `model.layers[18].mlp.up_proj.weight` and `model.layers[18].mlp.down_proj.weight`.

All the experiments are deployed on NVIDIA RTX 3090 Tensor Core GPUs, and we use 4 GPUs for training and single GPU for evaluation. For lifelong editing, due to computational constraints, we can accommodate a maximum of 5 experts. Consequently, the batch size in the sequence is determined by the total number of edits and the number of experts. For instance, when there are 100 edits and 5 experts, the batch size is set to 20; whereas with 1000 edits, the batch size scales up to 200. For bath editing in §6.1, the batch size is 30 and all the other parameters are the same as above.

B.4 Implementation for Ablation Study

In §6.3, we conducted an ablation study on several modules of LEMoE. Here, we detail the implementation of these ablations. In Table 4, Conventional routing means the router is modeled by a single-layer MLP, with the preservation of the insertion method. Knowledge routing is the knowledge anchor routing in MEMoE for short, also maintains the insertion method. Token-level embedding involves substituting e_j in Equation 12, which means $g(i | e_j) = g(i | x_i^{j_t})$. For entity-level embedding, we initially utilize the NLTK tool¹ for extracting named entities from the input instance. In cases where there are multiple named entities present in the input, we utilize the average pooling of the embeddings of these entities. Subsequently, we

¹<https://www.nltk.org>

replace e_j in Equation 10 with this embedding vector as the input to the sub-network to obtain the corresponding “value” of the input instances. All the other training hyperparameters are the same detailed in Appendix B.3.

C More Results and Analyses

C.1 More results for Influencing Factors

In §3.1, we employed two different evaluation methods: (1) a standard evaluation conducted on all edited data only after all edits were completed, and (2) an evaluation conducted immediately after each edit to assess the effectiveness of these edits at the current stage. Figure 2 shows the variations in the reliability metrics, and we further provide the changes in all three metrics here. To better illustrate these trends, we averaged the metrics over every four steps in a sequence of 100 editing steps. As shown in Figure 5, both reliability and generalization exhibit catastrophic forgetting phenomenon, where subsequent edits significantly affect the performance on prior data. This effect is most pronounced in the reliability metric. Additionally, around step 80, minimal fluctuations in the current editing reliability result in substantial oscillations in generalization. This can be attributed to the fact that, like human, a model must first accurately learn knowledge before it can generalize that knowledge. Thus, the generalization metric is, to some extent, contingent upon reliability. Regarding locality, the overall level remains consistently high, above 0.97, and thus the graph shows no discernible pattern of fluctuations. This further corroborates that knowledge editing through bypass mechanisms minimally impacts the model’s generalization capability (Wang and Li, 2024).

C.2 Case Study

In Table 7, we present bad cases of using LEMoE to edit the LLaMA-2-7B on the ZsRE dataset and mitigating these failures is critical for future work in model editing. We observe that:

i) errors occur only in part of the tokens, and these errors constitute a large proportion of the bad cases, indicating that the edits have not been sufficiently fitted. We wonder whether employing different learning rates and epochs for each batch in lifelong editing could alleviate this issue through more refined training.

ii) displays cases where the entire output is incorrect. These types of errors are the most common




Figure 5: Reliability, Generality and Locality of conventional MoE under different stage evaluation. “Immediate evaluation” occurs immediately after each edit, “Final evaluation” occurs after all edits in lifelong editing. Model : LLaMA2-7B. Dataset: ZsRE.

Table 7: Failure cases of LEMoE. ✕ represents errors in part of the tokens, ✕ represents complete output errors (i.e., factual failures), and ✓ indicates the expected exact match. Italics correspond to generalization prompt. LLaMA2-7B.

Prompt	Edit Target	Post-Edit Output
<i>i</i> What level is Javan surili’s iucn conservation status? <i>What is Javan surilis ucn conservation status?</i> The point in time of Air France Flight 447 was when? <i>When did Air France Flight 447 occur?</i>	critically threatened critically threatened 12 July 1944 12 July 1944	near threatened ✕ threatened ✕ 12 July 1967 ✕ 12 July 1967 ✕
<i>ii</i> Which war was William Babcock Hazen in? <i>What war did William Babcock Hazen go to?</i> When was the inception of Parcelforce? <i>When was Parcelforce formed?</i>	World War II World War II 1961 1961	US Civil War ✕ Spanish Civil War ✕ 1963 ✕ 1960 ✕
<i>iii</i> What team is Nicolas Raffault associated with? <i>Which team is Nicolas Raffault associated with?</i> What sports team was Petteri Nummelin a member of? <i>In which sports team was Petteri Nummelin a member?</i>	Arizona Coyotes Arizona Coyotes Columbus Blue Bombers Columbus Blue Bombers	Aqua ✕ Arizona Coyotes ✓ Cleveland Monsters ✕ Columbus Blue Bombers ✓
<i>iv</i> What level is Javan surili’s iucn conservation status? <i>What state is Qaleh Lan in?</i> When did Battle of the Java Sea occur? <i>When did the battle on the Java Sea begin?</i>	critically threatened critically threatened 27 February 1942 27 February 1942	nearly threatened ✓ unknown ✕ 27 February 1942 ✓ 1942 ✕

occurrences.

iv) presents cases of generalization failure. For example in prompt of last line, where the model answered “1942” which is partially correct, but did not fully follow the ground truth, indicating significant room for improvement in the accuracy of generalized edits.

Meanwhile, in *iii*) we surprisingly find that even when LEMoE errs on the Edit Prompt, it can correctly answer its paraphrase prompt. Upon closely examining these anomalous cases, we found that they predominantly pertain to question-answering scenarios within sports contexts, such as inquiries about a person’s team affiliation. We hypothesize that this phenomenon may stem from the relatively

limited number of teams in sports contexts, combined with the higher number of athletes and the occurrence of name duplication. Consequently, the model may accidentally provide correct answers to some of these questions.

In summary, LEMoE can handle contextual information correctly in some cases but falls short in specific editing instructions, suggesting that optimizing editing instructions (modifying the editing context) may be a direction for improvement.

C.3 More Ablation Results

As an extension §6.3, we evaluate the effectiveness of LEMoE applied to different layers. For experimental setup, we utilize the LLama2-7b model and

Table 8: Performance of LEMoE on different layer of LLaMA2-7B using ZsRE.

Layer	Rel.↑	Gen.↑	Loc.↑	Avg.↑
0	0.28	0.14	1.00	0.47
2	0.55	0.41	1.00	0.65
4	0.52	0.34	1.00	0.62
6	0.44	0.25	1.00	0.56
8	0.48	0.26	1.00	0.58
10	0.47	0.25	1.00	0.57
12	0.54	0.26	1.00	0.60
14	0.58	0.33	1.00	0.64
16	0.77	0.55	1.00	0.77
18	0.80	0.60	1.00	0.80
20	0.76	0.59	1.00	0.79
22	0.70	0.54	1.00	0.75
24	0.74	0.51	1.00	0.75
26	0.77	0.56	1.00	0.78
28	0.73	0.48	1.00	0.74
30	0.43	0.24	1.00	0.56

the ZsRE dataset. Lifelong editing involves 1000 instances, with all the other training hyperparameters consistent as detailed in Appendix B.3. Experimental results are depicted in Table 8. Notably, the 18-th layer exhibits the most significant editing improvements, achieving peak performance across all metrics. Conversely, the first layer demonstrates the least improvement, and the editing hardly takes effect in the low-level transformer block. In contrast, high-level transformer blocks display pronounced editing effects, maintaining high reliability consistently from the 16-th layer onwards. However, significant degradation in editing efficacy is noted towards the 30-th layer, possibly due to the increased proximity to the output. On the other hand, locality remains unaffected, consistently scoring 1.00. Thus, our findings further validate that the high-level transformer blocks of LM based on the transformer architecture contain factual information, and editing of these layers will have a significant effect (Yao et al., 2024).

C.4 LEMoE with LoRA structure

In the era of LLMs, parameter-efficient fine-tuning (PEFT) methods such as LoRA have proven highly effective and convenient for achieving impressive results across various downstream tasks. LoRA (Hu et al., 2022), proposes a technique that decomposes the update gradient matrix into two small rank- n matrices, significantly reducing the memory

requirements for training LLMs. Meanwhile, in fields of MoE, some studies have explored replacing traditional MoE structures with LoRA (Zadouri et al., 2023; Wu et al., 2024b). Consequently, we replace the MLP-based expert networks in LEMoE with LoRA modules. Given the challenging nature of lifelong learning tasks, we evaluate the performance of this low-parameter model structure on batch editing tasks with batch size set to 30.

We investigated the effects of varying the number of experts (Exp.), different LoRA ranks, and different top_k values. Detailed experimental results are provided in the Table 9 to facilitate further research. We conducted experiments on all even-numbered layers, expert number in [1,10,20], top_k in [1,10,20] and LoRA Rank in [2,4,8,16,32,64,128,256,512,1024,2048]. We filter out results with Reliability below 0.1, Generalization below 0.1, and Locality below 0.5. It is evident that this method performs poorly in editing the low-level transformer blocks (with results falling below the selection criteria and many being zero, hence not presented in the table). Meanwhile, the higher the layer being edited, the better the performance observed. This LEMoE-LoRA achieved optimal performance with 30 layers, 10 experts, a LoRA rank of 2048, and a top_k value of 10.

Table 9: Experimental Results of LEMoE with LoRA module.

Layer	Expert Number	Topk	LoRA Rank	Rel. $\uparrow$	Gen. $\uparrow$	Loc. $\uparrow$	Avg. $\uparrow$
4	10	10	1024	0.13	0.10	0.53	0.26
4	10	10	2048	0.30	0.23	0.68	0.41
4	20	5	512	0.23	0.17	0.58	0.33
4	20	10	1024	0.30	0.27	0.50	0.36
6	20	5	1024	0.17	0.17	0.67	0.33
8	1	1	1024	0.13	0.10	0.62	0.28
8	1	1	2048	0.30	0.20	0.90	0.47
8	10	1	512	0.17	0.17	0.65	0.33
8	10	5	2048	0.40	0.13	0.67	0.40
8	10	10	1024	0.13	0.10	0.62	0.28
8	10	10	2048	0.23	0.10	0.78	0.37
8	20	5	1024	0.10	0.10	0.77	0.32
8	20	10	512	0.17	0.13	0.90	0.40
10	10	5	2048	0.13	0.20	0.93	0.42
10	20	5	512	0.13	0.13	0.93	0.40
10	20	5	1024	0.23	0.17	0.93	0.44
10	20	10	1024	0.10	0.10	0.95	0.38
12	1	1	2048	0.30	0.13	0.98	0.47
12	10	1	1024	0.17	0.10	0.68	0.32
12	10	5	1024	0.27	0.17	0.95	0.46
12	10	5	2048	0.23	0.13	0.98	0.45
12	10	10	1024	0.23	0.10	0.93	0.42
12	10	10	2048	0.23	0.17	0.98	0.46
12	20	5	512	0.20	0.10	0.98	0.43
12	20	5	1024	0.37	0.23	0.97	0.52
12	20	10	512	0.23	0.13	0.98	0.45
12	20	10	1024	0.20	0.10	0.98	0.43
14	1	1	16	0.13	0.10	0.97	0.40
14	1	1	512	0.17	0.10	0.98	0.42
14	1	1	1024	0.30	0.10	0.98	0.46
14	1	1	2048	0.33	0.17	1.00	0.50
14	10	1	128	0.13	0.10	0.68	0.31
14	10	1	512	0.40	0.30	0.98	0.56
14	10	1	1024	0.40	0.23	1.00	0.54
14	10	1	2048	0.57	0.43	1.00	0.67
14	10	5	512	0.27	0.17	0.98	0.47
14	10	5	1024	0.23	0.20	0.98	0.47
14	10	5	2048	0.50	0.43	1.00	0.64
14	10	10	16	0.13	0.10	0.95	0.39
14	10	10	128	0.20	0.20	1.00	0.47
14	10	10	512	0.23	0.20	1.00	0.48
14	10	10	1024	0.27	0.13	1.00	0.47
14	10	10	2048	0.40	0.17	1.00	0.52
14	20	1	1024	0.37	0.37	0.88	0.54
14	20	5	128	0.17	0.13	0.98	0.43
14	20	5	512	0.37	0.20	1.00	0.52
14	20	5	1024	0.40	0.23	0.97	0.53

Continued on next page

Table 9 Continued from previous page

Layer	Expert Number	Topk	LoRA Rank	Rel. ↑	Gen. ↑	Loc. ↑	Avg. ↑
14	20	10	128	0.10	0.10	0.98	0.39
14	20	10	512	0.20	0.13	1.00	0.44
14	20	10	1024	0.30	0.20	1.00	0.50
16	1	1	128	0.20	0.10	1.00	0.43
16	1	1	512	0.37	0.30	0.98	0.55
16	1	1	1024	0.60	0.43	1.00	0.68
16	1	1	2048	0.60	0.43	0.98	0.67
16	10	1	512	0.43	0.37	0.95	0.58
16	10	1	1024	0.47	0.33	0.93	0.58
16	10	1	2048	0.73	0.50	0.93	0.72
16	10	5	16	0.17	0.10	0.97	0.41
16	10	5	128	0.17	0.13	1.00	0.43
16	10	5	512	0.50	0.27	1.00	0.59
16	10	5	1024	0.37	0.17	1.00	0.51
16	10	5	2048	0.53	0.20	1.00	0.58
16	10	10	16	0.20	0.10	1.00	0.43
16	10	10	128	0.33	0.27	1.00	0.53
16	10	10	512	0.40	0.33	1.00	0.58
16	10	10	1024	0.57	0.43	0.97	0.66
16	10	10	2048	0.63	0.33	0.98	0.65
16	20	1	128	0.33	0.13	0.80	0.42
16	20	1	512	0.37	0.17	0.97	0.50
16	20	1	1024	0.40	0.33	1.00	0.58
16	20	5	128	0.33	0.27	1.00	0.53
16	20	5	512	0.47	0.27	1.00	0.58
16	20	5	1024	0.70	0.43	1.00	0.71
16	20	10	128	0.17	0.13	1.00	0.43
16	20	10	512	0.23	0.10	1.00	0.44
16	20	10	1024	0.43	0.23	1.00	0.56
18	1	1	128	0.43	0.33	1.00	0.59
18	1	1	512	0.50	0.33	1.00	0.61
18	1	1	1024	0.57	0.37	1.00	0.64
18	1	1	2048	0.77	0.53	1.00	0.77
18	10	1	128	0.23	0.13	0.93	0.43
18	10	1	512	0.53	0.33	0.98	0.62
18	10	1	1024	0.20	0.20	0.77	0.39
18	10	1	2048	0.33	0.30	0.95	0.53
18	10	5	16	0.30	0.13	0.98	0.47
18	10	5	128	0.30	0.23	1.00	0.51
18	10	5	512	0.43	0.23	1.00	0.56
18	10	5	1024	0.53	0.40	1.00	0.64
18	10	5	2048	0.73	0.53	1.00	0.76
18	10	10	16	0.17	0.13	1.00	0.43
18	10	10	128	0.47	0.30	1.00	0.59
18	10	10	512	0.63	0.37	1.00	0.67
18	10	10	1024	0.67	0.40	1.00	0.69
18	10	10	2048	0.70	0.53	1.00	0.74
18	20	1	128	0.30	0.20	0.92	0.47

Continued on next page

Table 9 Continued from previous page

Layer	Expert Number	Topk	LoRA Rank	Rel. $\uparrow$	Gen. $\uparrow$	Loc. $\uparrow$	Avg. $\uparrow$
18	20	1	512	0.40	0.27	0.90	0.52
18	20	1	1024	0.47	0.43	0.98	0.63
18	20	5	16	0.27	0.17	0.75	0.39
18	20	5	128	0.47	0.27	0.98	0.57
18	20	5	512	0.53	0.33	0.98	0.62
18	20	5	1024	0.67	0.50	1.00	0.72
18	20	10	16	0.27	0.17	0.97	0.47
18	20	10	128	0.30	0.27	1.00	0.52
18	20	10	512	0.50	0.33	1.00	0.61
18	20	10	1024	0.57	0.30	1.00	0.62
20	1	1	128	0.50	0.27	0.98	0.58
20	1	1	512	0.60	0.33	1.00	0.64
20	1	1	1024	0.63	0.33	1.00	0.66
20	1	1	2048	0.70	0.40	1.00	0.70
20	10	1	128	0.47	0.37	0.93	0.59
20	10	1	512	0.40	0.43	0.97	0.60
20	10	1	1024	0.33	0.20	0.88	0.47
20	10	1	2048	0.37	0.37	0.85	0.53
20	10	5	16	0.37	0.20	1.00	0.52
20	10	5	128	0.50	0.20	1.00	0.57
20	10	5	512	0.40	0.23	1.00	0.54
20	10	5	1024	0.57	0.23	1.00	0.60
20	10	5	2048	0.67	0.30	1.00	0.66
20	10	10	16	0.30	0.20	1.00	0.50
20	10	10	128	0.47	0.20	1.00	0.56
20	10	10	512	0.37	0.23	1.00	0.53
20	10	10	1024	0.60	0.30	1.00	0.63
20	10	10	2048	0.60	0.40	1.00	0.67
20	20	1	128	0.20	0.20	0.68	0.36
20	20	1	512	0.33	0.30	1.00	0.54
20	20	1	1024	0.57	0.33	0.92	0.61
20	20	5	16	0.43	0.30	0.90	0.54
20	20	5	128	0.47	0.20	1.00	0.56
20	20	5	512	0.47	0.30	1.00	0.59
20	20	5	1024	0.60	0.43	0.97	0.67
20	20	10	16	0.33	0.23	0.97	0.51
20	20	10	128	0.47	0.27	1.00	0.58
20	20	10	512	0.53	0.23	1.00	0.59
20	20	10	1024	0.50	0.30	1.00	0.60
22	1	1	128	0.40	0.17	0.98	0.52
22	1	1	512	0.40	0.23	1.00	0.54
22	1	1	1024	0.43	0.30	1.00	0.58
22	1	1	2048	0.50	0.30	1.00	0.60
22	10	1	128	0.23	0.17	0.97	0.46
22	10	1	512	0.23	0.33	0.95	0.51
22	10	1	2048	0.40	0.17	0.90	0.49
22	10	5	16	0.23	0.20	1.00	0.48
22	10	5	128	0.33	0.23	1.00	0.52

Continued on next page

Table 9 Continued from previous page

Layer	Expert Number	Topk	LoRA Rank	Rel. ↑	Gen. ↑	Loc. ↑	Avg. ↑
22	10	5	512	0.47	0.33	1.00	0.60
22	10	5	1024	0.47	0.20	1.00	0.56
22	10	5	2048	0.50	0.27	1.00	0.59
22	10	10	16	0.20	0.13	1.00	0.44
22	10	10	128	0.40	0.17	1.00	0.52
22	10	10	512	0.40	0.20	1.00	0.53
22	10	10	1024	0.40	0.23	1.00	0.54
22	10	10	2048	0.43	0.30	1.00	0.58
22	20	1	128	0.23	0.17	0.95	0.45
22	20	1	512	0.17	0.17	1.00	0.44
22	20	1	1024	0.40	0.30	0.87	0.52
22	20	5	16	0.30	0.13	1.00	0.48
22	20	5	128	0.37	0.20	1.00	0.52
22	20	5	512	0.47	0.20	0.97	0.54
22	20	5	1024	0.47	0.23	1.00	0.57
22	20	10	16	0.33	0.13	1.00	0.49
22	20	10	128	0.43	0.23	1.00	0.56
22	20	10	512	0.43	0.30	1.00	0.58
22	20	10	1024	0.47	0.23	1.00	0.57
24	1	1	16	0.20	0.17	1.00	0.46
24	1	1	128	0.40	0.30	1.00	0.57
24	1	1	512	0.43	0.30	1.00	0.58
24	1	1	1024	0.50	0.33	1.00	0.61
24	1	1	2048	0.57	0.47	0.98	0.67
24	10	1	512	0.37	0.20	0.97	0.51
24	10	1	1024	0.20	0.10	0.93	0.41
24	10	1	2048	0.27	0.27	0.93	0.49
24	10	5	16	0.33	0.17	1.00	0.50
24	10	5	128	0.50	0.40	1.00	0.63
24	10	5	512	0.60	0.47	1.00	0.69
24	10	5	1024	0.57	0.47	1.00	0.68
24	10	5	2048	0.53	0.37	1.00	0.63
24	10	10	16	0.43	0.23	1.00	0.56
24	10	10	128	0.50	0.43	1.00	0.64
24	10	10	512	0.37	0.40	1.00	0.59
24	10	10	1024	0.57	0.50	1.00	0.69
24	10	10	2048	0.60	0.47	1.00	0.69
24	20	1	128	0.20	0.20	0.93	0.44
24	20	1	512	0.33	0.33	1.00	0.56
24	20	5	16	0.30	0.23	1.00	0.51
24	20	5	128	0.53	0.43	1.00	0.66
24	20	5	512	0.53	0.33	0.98	0.62
24	20	5	1024	0.63	0.50	1.00	0.71
24	20	10	16	0.37	0.20	0.98	0.52
24	20	10	128	0.57	0.40	1.00	0.66
24	20	10	512	0.53	0.47	1.00	0.67
24	20	10	1024	0.47	0.40	1.00	0.62
26	1	1	16	0.13	0.13	1.00	0.42

Continued on next page

Table 9 Continued from previous page

Layer	Expert Number	Topk	LoRA Rank	Rel. ↑	Gen. ↑	Loc. ↑	Avg. ↑
26	1	1	128	0.57	0.37	1.00	0.64
26	1	1	512	0.53	0.50	1.00	0.68
26	1	1	1024	0.53	0.47	1.00	0.67
26	1	1	2048	0.63	0.50	1.00	0.71
26	10	1	512	0.27	0.23	0.95	0.48
26	10	1	1024	0.13	0.10	0.93	0.39
26	10	1	2048	0.37	0.33	0.95	0.55
26	10	5	2	0.10	0.10	0.90	0.37
26	10	5	16	0.50	0.30	1.00	0.60
26	10	5	128	0.57	0.37	1.00	0.64
26	10	5	512	0.57	0.43	1.00	0.67
26	10	5	1024	0.63	0.53	1.00	0.72
26	10	5	2048	0.50	0.57	1.00	0.69
26	10	10	16	0.47	0.30	1.00	0.59
26	10	10	128	0.53	0.37	1.00	0.63
26	10	10	512	0.53	0.43	1.00	0.66
26	10	10	1024	0.53	0.43	1.00	0.66
26	10	10	2048	0.63	0.57	0.98	0.73
26	20	1	16	0.10	0.13	0.72	0.32
26	20	1	512	0.23	0.20	0.90	0.44
26	20	1	1024	0.40	0.33	0.95	0.56
26	20	5	16	0.53	0.33	1.00	0.62
26	20	5	128	0.70	0.47	1.00	0.72
26	20	5	512	0.60	0.47	1.00	0.69
26	20	5	1024	0.53	0.30	1.00	0.61
26	20	10	2	0.17	0.10	0.80	0.36
26	20	10	16	0.53	0.43	1.00	0.66
26	20	10	128	0.57	0.30	1.00	0.62
26	20	10	512	0.60	0.53	1.00	0.71
26	20	10	1024	0.57	0.40	1.00	0.66
28	1	1	16	0.10	0.13	0.83	0.36
28	1	1	128	0.47	0.33	0.93	0.58
28	1	1	512	0.50	0.43	0.97	0.63
28	1	1	1024	0.60	0.37	0.98	0.65
28	1	1	2048	0.60	0.50	1.00	0.70
28	10	1	128	0.30	0.23	0.70	0.41
28	10	1	1024	0.30	0.10	0.87	0.42
28	10	1	2048	0.30	0.27	0.88	0.48
28	10	5	16	0.37	0.40	1.00	0.59
28	10	5	128	0.47	0.37	0.83	0.56
28	10	5	512	0.47	0.47	1.00	0.64
28	10	5	1024	0.70	0.60	1.00	0.77
28	10	5	2048	0.70	0.47	1.00	0.72
28	10	10	16	0.30	0.30	0.98	0.53
28	10	10	128	0.50	0.33	0.95	0.59
28	10	10	512	0.63	0.40	0.97	0.67
28	10	10	1024	0.70	0.40	1.00	0.70
28	10	10	2048	0.60	0.40	1.00	0.67

Continued on next page

Table 9 Continued from previous page

Layer	Expert Number	Topk	LoRA Rank	Rel. ↑	Gen. ↑	Loc. ↑	Avg. ↑
28	20	1	512	0.27	0.23	0.95	0.48
28	20	1	1024	0.27	0.10	0.93	0.43
28	20	5	16	0.47	0.33	0.98	0.59
28	20	5	128	0.50	0.40	0.97	0.62
28	20	5	512	0.50	0.37	0.92	0.59
28	20	5	1024	0.50	0.37	1.00	0.62
28	20	10	16	0.43	0.27	0.98	0.56
28	20	10	128	0.43	0.33	0.98	0.58
28	20	10	512	0.60	0.40	1.00	0.67
28	20	10	1024	0.50	0.30	1.00	0.60
30	1	1	16	0.23	0.17	1.00	0.47
30	1	1	128	0.67	0.47	1.00	0.71
30	1	1	512	0.73	0.57	1.00	0.77
30	1	1	1024	0.70	0.63	1.00	0.78
30	1	1	2048	0.73	0.73	1.00	0.82
30	10	1	128	0.23	0.10	0.95	0.43
30	10	1	512	0.20	0.20	0.95	0.45
30	10	5	16	0.53	0.40	1.00	0.64
30	10	5	128	0.57	0.63	1.00	0.73
30	10	5	512	0.53	0.63	1.00	0.72
30	10	5	1024	0.70	0.63	1.00	0.78
30	10	5	2048	0.67	0.50	1.00	0.72
30	10	10	16	0.43	0.27	1.00	0.57
30	10	10	128	0.57	0.57	1.00	0.71
30	10	10	512	0.60	0.50	1.00	0.70
30	10	10	1024	0.73	0.63	1.00	0.79
30	10	10	2048	0.77	0.73	1.00	0.83
30	20	1	16	0.10	0.10	0.63	0.28
30	20	1	128	0.17	0.13	1.00	0.43
30	20	1	512	0.27	0.27	1.00	0.51
30	20	1	1024	0.20	0.13	0.93	0.42
30	20	5	16	0.47	0.37	1.00	0.61
30	20	5	128	0.53	0.57	1.00	0.70
30	20	5	512	0.60	0.40	1.00	0.67
30	20	5	1024	0.43	0.37	1.00	0.60
30	20	10	16	0.53	0.43	1.00	0.66
30	20	10	128	0.53	0.43	1.00	0.66
30	20	10	512	0.63	0.50	1.00	0.71
30	20	10	1024	0.47	0.37	1.00	0.61